

Group Name: " "

Member 1: Rodel Yzer C. De Boda

SR Code: 22-04506

Member 2: Adhem V. Hernandez

SR Code: 22-07415

✓ Laboratory Activity 2: Applied Classification Analysis

In this laboratory we will create an classification algorithm for our gathered classification datasets.

Dataset Link: <https://www.kaggle.com/datasets/erdemtaha/cancer-data/data>

✓ Description of the dataset

Insert Here: "The dataset contains the characteristics of patients diagnosed with cancer. The dataset contains a unique ID for each patient, the type of cancer (diagnosis), the visual characteristics of the cancer and the average values of these characteristics."

✓ Import the Required Packages

For this exercise we will require the Pandas package for loading the data, the matplotlib package for plotting as well as scikit-learn for creating the Classification model. Import all of the required packages and relevant modules for these tasks.

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
import missingno as msno

import numpy as np
from array import array
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.metrics import accuracy_score, f1_score, precision_score
```

✓ Read the Data

```
df = pd.read_csv("/content/Cancer_Data.csv")
```

✓ Read the data and find summary statistics

Get the info of your dataset


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    569 non-null   int64
1   diagnosis             569 non-null   object
2   radius_mean           569 non-null   float64
3   texture_mean          569 non-null   float64
4   perimeter_mean        569 non-null   float64
5   area_mean             569 non-null   float64
```

```
6 smoothness_mean      569 non-null    float64
7 compactness_mean     569 non-null    float64
8 concavity_mean       569 non-null    float64
9 concave points_mean  569 non-null    float64
10 symmetry_mean       569 non-null    float64
11 fractal_dimension_mean 569 non-null    float64
12 radius_se          569 non-null    float64
13 texture_se         569 non-null    float64
14 perimeter_se       569 non-null    float64
15 area_se            569 non-null    float64
16 smoothness_se      569 non-null    float64
17 compactness_se     569 non-null    float64
18 concavity_se       569 non-null    float64
19 concave points_se  569 non-null    float64
20 symmetry_se        569 non-null    float64
21 fractal_dimension_se 569 non-null    float64
22 radius_worst       569 non-null    float64
23 texture_worst      569 non-null    float64
24 perimeter_worst    569 non-null    float64
25 area_worst         569 non-null    float64
26 smoothness_worst   569 non-null    float64
27 compactness_worst  569 non-null    float64
28 concavity_worst    569 non-null    float64
29 concave points_worst 569 non-null    float64
30 symmetry_worst     569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
32 Unnamed: 32        0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

Get the first five and last five of your dataset


```
df.head()
```



| | id | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean |
|---|----------|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|
| 0 | 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.1471 |
| 1 | 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.0974 |
| 2 | 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.1279 |
| 3 | 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.1172 |
| 4 | 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.1154 |

5 rows × 33 columns

```
df.tail(5)
```



| | id | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean |
|-----|--------|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|
| 564 | 926424 | M | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | 0.11590 | 0.24390 | 0.1154 |
| 565 | 926682 | M | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | 0.10340 | 0.14400 | 0.0974 |
| 566 | 926954 | M | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | 0.10230 | 0.09251 | 0.0974 |
| 567 | 927241 | M | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | 0.27700 | 0.35140 | 0.1154 |
| 568 | 92751 | B | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | 0.04362 | 0.00000 | 0.0974 |

5 rows × 33 columns

Get the summary statistics of your dataset to show the total count, mean, standard deviation, min and max value, and percentiles of each column of your dataset.

```
df.describe()
```



| | id | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | conca points_me |
|-------|--------------|-------------|--------------|----------------|-------------|-----------------|------------------|----------------|--------------------|
| count | 5.690000e+02 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.0000 |
| mean | 3.037183e+07 | 14.127292 | 19.289649 | 91.969033 | 654.889104 | 0.096360 | 0.104341 | 0.088799 | 0.0489 |
| std | 1.250206e+08 | 3.524049 | 4.301036 | 24.298981 | 351.914129 | 0.014064 | 0.052813 | 0.079720 | 0.0388 |
| min | 8.670000e+03 | 6.981000 | 9.710000 | 43.790000 | 143.500000 | 0.052630 | 0.019380 | 0.000000 | 0.0000 |
| 25% | 8.692180e+05 | 11.700000 | 16.170000 | 75.170000 | 420.300000 | 0.086370 | 0.064920 | 0.029560 | 0.0203 |
| 50% | 9.060240e+05 | 13.370000 | 18.840000 | 86.240000 | 551.100000 | 0.095870 | 0.092630 | 0.061540 | 0.0335 |
| 75% | 8.813129e+06 | 15.780000 | 21.800000 | 104.100000 | 782.700000 | 0.105300 | 0.130400 | 0.130700 | 0.0740 |
| max | 9.113205e+08 | 28.110000 | 39.280000 | 188.500000 | 2501.000000 | 0.163400 | 0.345400 | 0.426800 | 0.2012 |

8 rows × 32 columns

▼ Data Cleaning

Visualize the missing values/data in your dataset to see how many data and percentage of data are missing in each column of your dataset.


```
df.head()
```



| | id | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | conci points_i |
|---|----------|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|-------------------|
| 0 | 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14 |
| 1 | 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07 |
| 2 | 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.17 |
| 3 | 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10 |
| 4 | 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10 |

5 rows × 33 columns

```
df.isnull().sum()
```

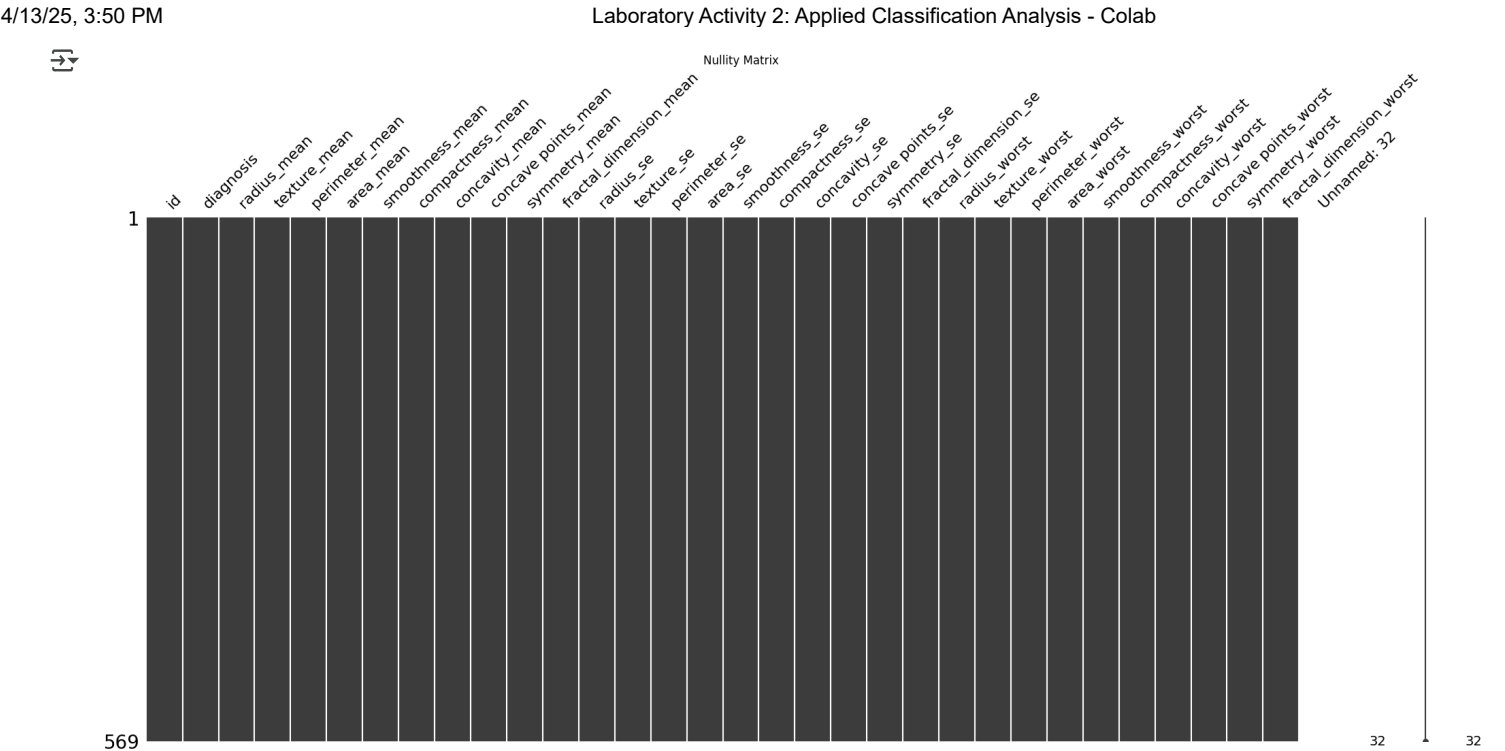


| | 0 |
|-------------------------|-----|
| id | 0 |
| diagnosis | 0 |
| radius_mean | 0 |
| texture_mean | 0 |
| perimeter_mean | 0 |
| area_mean | 0 |
| smoothness_mean | 0 |
| compactness_mean | 0 |
| concavity_mean | 0 |
| concave points_mean | 0 |
| symmetry_mean | 0 |
| fractal_dimension_mean | 0 |
| radius_se | 0 |
| texture_se | 0 |
| perimeter_se | 0 |
| area_se | 0 |
| smoothness_se | 0 |
| compactness_se | 0 |
| concavity_se | 0 |
| concave points_se | 0 |
| symmetry_se | 0 |
| fractal_dimension_se | 0 |
| radius_worst | 0 |
| texture_worst | 0 |
| perimeter_worst | 0 |
| area_worst | 0 |
| smoothness_worst | 0 |
| compactness_worst | 0 |
| concavity_worst | 0 |
| concave points_worst | 0 |
| symmetry_worst | 0 |
| fractal_dimension_worst | 0 |
| Unnamed: 32 | 569 |

df = df

Use a nullity matrix for your dataset for easy visualization of missing data

```
msno.matrix(df)
plt.title('Nullity Matrix')
plt.show()
```




Imputation of data to missing data

```
#drop data with 80% missing percentage value
columns_to_drop = df.columns[df.isnull().sum() / len(df) * 100 > 80]
df.drop(columns_to_drop, axis=1, inplace=True)
df.head()
```

| | id | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean |
|---|----------|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|
| 0 | 842302 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14 |
| 1 | 842517 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07 |
| 2 | 84300903 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12 |
| 3 | 84348301 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10 |
| 4 | 84358402 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10 |

5 rows × 32 columns

```
#dropping or deleting columns that is unnecessary
df.drop(['id'], axis=1, inplace=True)
df.head()
```




| | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | symmetry_mean |
|---|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|---------------|
| 0 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | |
| 1 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | |
| 2 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | |
| 3 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | |
| 4 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | |

5 rows × 31 columns

After deleting the columns with 80% missing data. Impute data/value in the missing data of your dataset.

```
features = df.columns.drop(['diagnosis'])
for column in features:
    df[column].fillna(df[column].mean(), inplace=True)
```




<ipython-input-155-ea25aba2fed0>:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assign
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting valu

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].me

```
df[column].fillna(df[column].mean(), inplace=True)
```

After imputation of data, Check the completeness of your dataset by getting its info.

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    object
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                       569 non-null    float64
6   compactness_mean                      569 non-null    float64
7   concavity_mean                        569 non-null    float64
8   concave points_mean                   569 non-null    float64
9   symmetry_mean                         569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                             569 non-null    float64
12  texture_se                             569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                               569 non-null    float64
15  smoothness_se                         569 non-null    float64
16  compactness_se                        569 non-null    float64
17  concavity_se                          569 non-null    float64
18  concave points_se                     569 non-null    float64
19  symmetry_se                           569 non-null    float64
20  fractal_dimension_se                   569 non-null    float64
21  radius_worst                          569 non-null    float64
22  texture_worst                         569 non-null    float64
23  perimeter_worst                       569 non-null    float64
24  area_worst                            569 non-null    float64
25  smoothness_worst                      569 non-null    float64
26  compactness_worst                     569 non-null    float64
27  concavity_worst                       569 non-null    float64
28  concave points_worst                   569 non-null    float64
29  symmetry_worst                        569 non-null    float64
30  fractal_dimension_worst                569 non-null    float64
dtypes: float64(30), object(1)
memory usage: 137.9+ KB
```

Visualize the Data

Load the dataset using Pandas and plot the different the target category

Plot a number of different features vs the allocated species classifications e.g. Sepal Length vs Petal length and Species. Visually inspect the plots and look for any patterns that could indicate separation between each of the species.

```
df['diagnosis'].unique()

array(['M', 'B'], dtype=object)
```

✓ Preprocessing

Feature Engineering

You need to select the most appropriate features that will provide the most powerful classification model.

Before we can construct the model we must first convert the species values into labels that can be used within the model. Replace:

Example:

1. The species string Iris-setosa with the value 0
2. The species string Iris-versicolor with the value 1
3. The species string Iris-virginica with the value 2



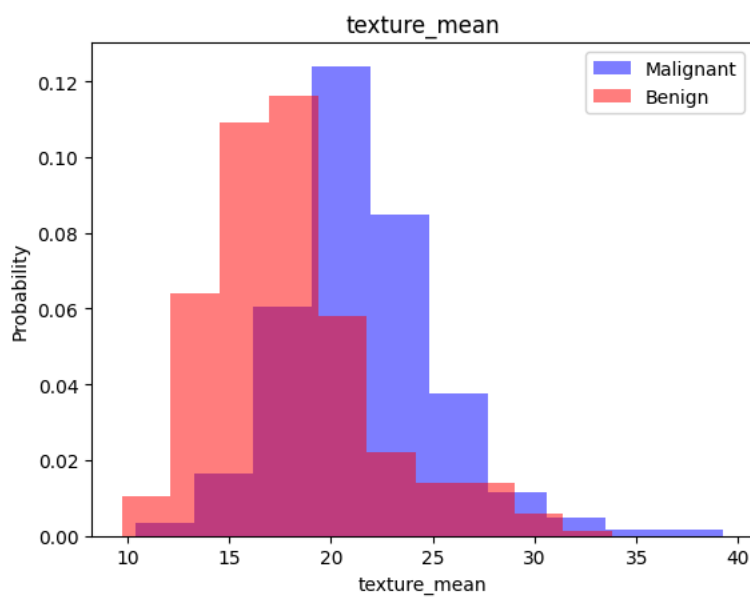
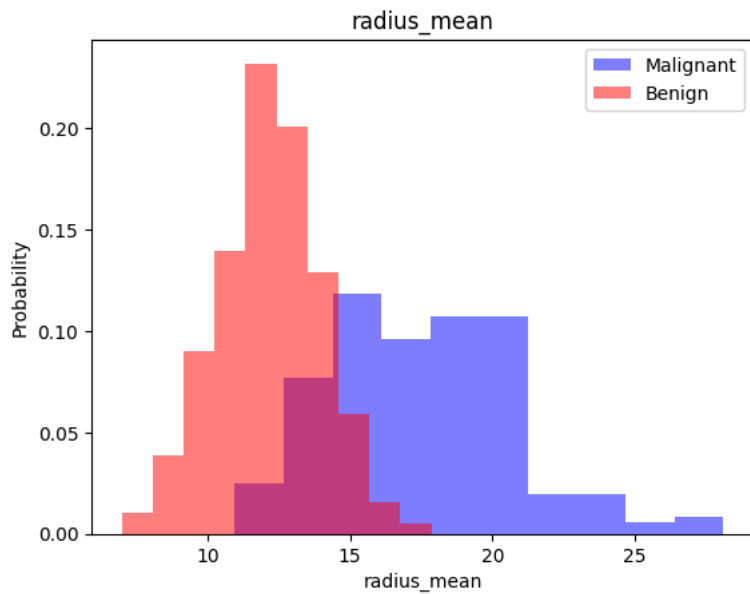
You can use get_dummies function / OnehotEncoder / LabelEncoder Library / StandardScaler with this.


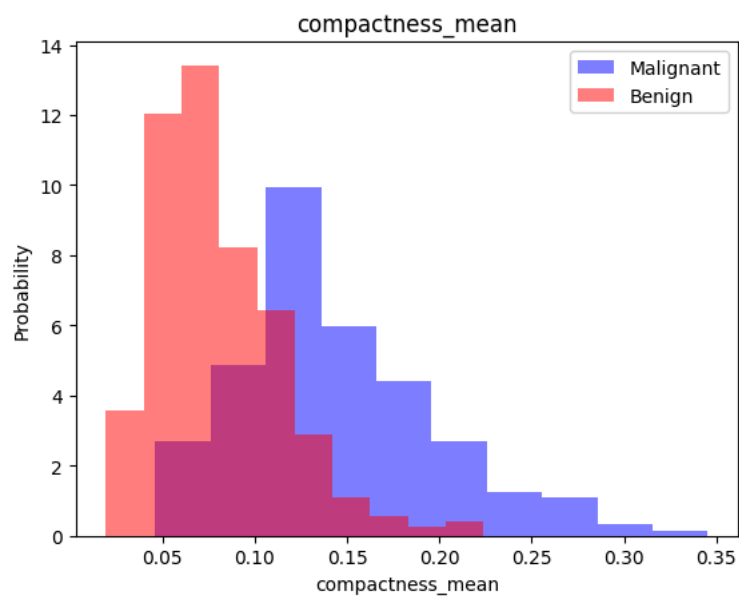


features

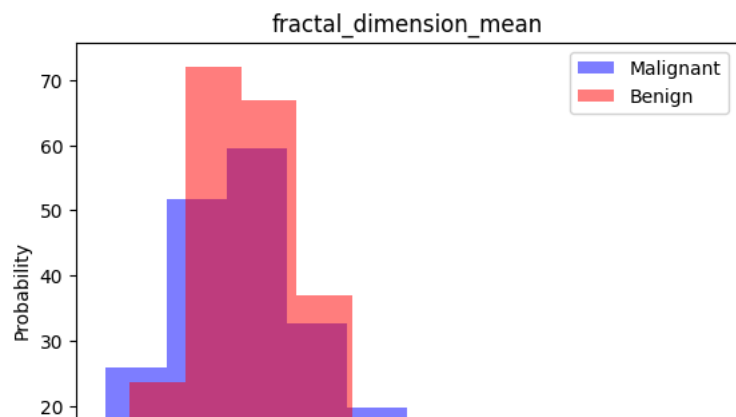
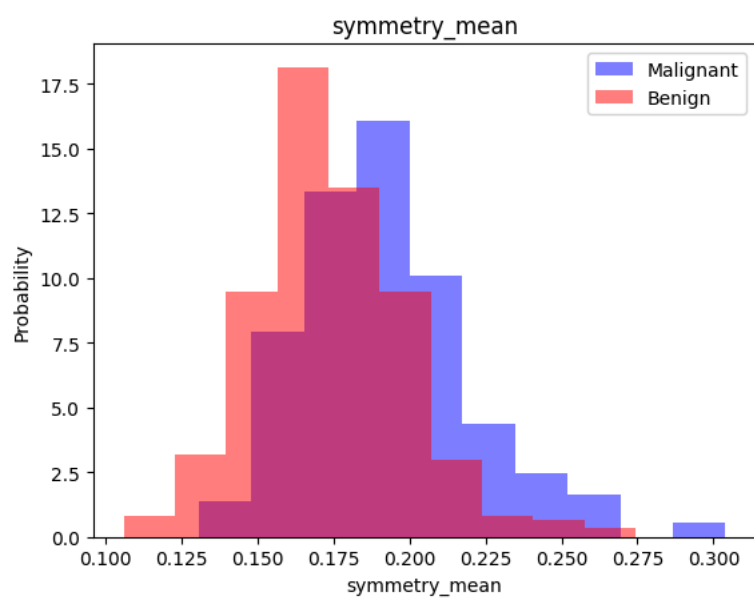

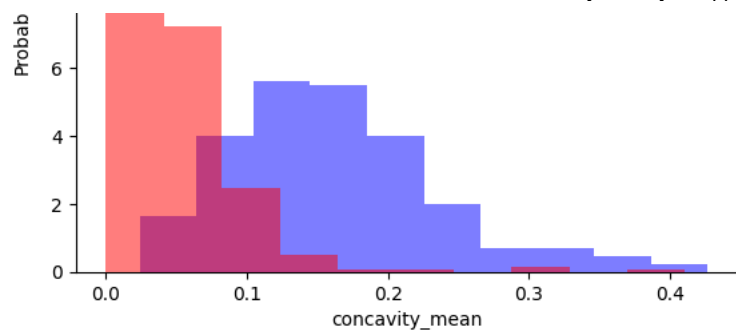
```
Index(['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean',
       'smoothness_mean', 'compactness_mean', 'concavity_mean',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst'],
      dtype='object')

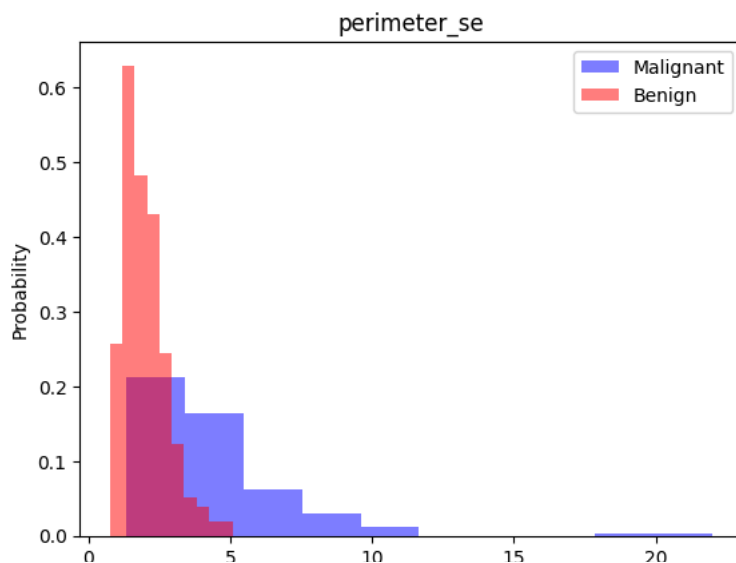



df['diagnosis'] = (df['diagnosis'] == 'M').astype(int)



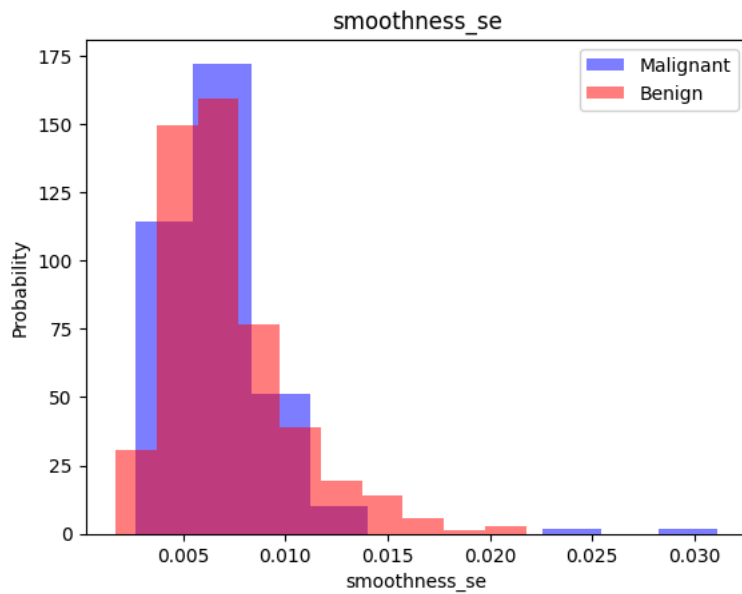
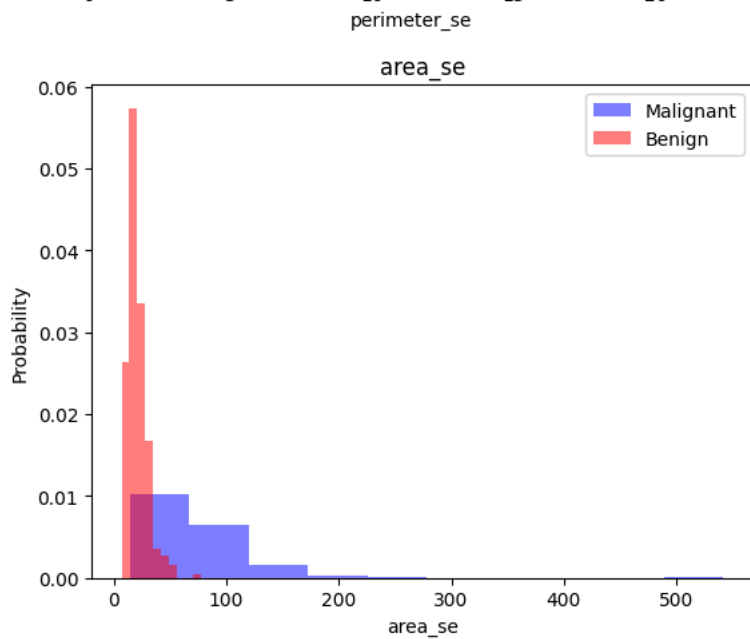
for feature in features:
    plt.hist(df[df['diagnosis']==1][feature], color='blue', label='Malignant', density=True, alpha=0.5)
    plt.hist(df[df['diagnosis']==0][feature], color='red', label='Benign', density=True, alpha=0.5)
    plt.ylabel('Probability')
    plt.xlabel(feature)
    plt.legend()
    plt.title(feature)
    plt.show()
```



































```
# Get a list of all columns
cols = df.columns.tolist()

# Move the first column to the end
cols = cols[1:] + cols[:1]

# Reorder the DataFrame
df = df[cols]

train, valid, test = np.split(df.sample(frac=1), [int(0.6 * len(df)), int(0.8 * len(df))])

/usr/local/lib/python3.11/dist-packages/numpy/_core/fromnumeric.py:57: FutureWarning: 'DataFrame.swapaxes' is deprecated and will be removed in a future version.
  return bound(*args, **kwargs)
```

train

| | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | symmetry_mean |
|-----|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|---------------|
| 266 | 10.600 | 18.95 | 69.28 | 346.4 | 0.09688 | 0.11470 | 0.06387 | 0.02642 | 0.1922 |
| 321 | 20.160 | 19.66 | 131.10 | 1274.0 | 0.08020 | 0.08564 | 0.11550 | 0.07726 | 0.1928 |
| 511 | 14.810 | 14.70 | 94.66 | 680.7 | 0.08472 | 0.05016 | 0.03416 | 0.02541 | 0.1659 |
| 491 | 17.850 | 13.23 | 114.60 | 992.1 | 0.07838 | 0.06217 | 0.04445 | 0.04178 | 0.1220 |
| 318 | 9.042 | 18.90 | 60.07 | 244.5 | 0.09968 | 0.19720 | 0.19750 | 0.04908 | 0.2330 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 422 | 11.610 | 16.02 | 75.46 | 408.2 | 0.10880 | 0.11680 | 0.07097 | 0.04497 | 0.1886 |
| 31 | 11.840 | 18.70 | 77.93 | 440.6 | 0.11090 | 0.15160 | 0.12180 | 0.05182 | 0.2301 |
| 26 | 14.580 | 21.53 | 97.41 | 644.8 | 0.10540 | 0.18680 | 0.14250 | 0.08783 | 0.2252 |
| 217 | 10.200 | 17.48 | 65.05 | 321.2 | 0.08054 | 0.05907 | 0.05774 | 0.01071 | 0.1964 |
| 478 | 11.490 | 14.59 | 73.99 | 404.9 | 0.10460 | 0.08228 | 0.05308 | 0.01969 | 0.1779 |

341 rows × 31 columns

test

| | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | symmetry_mean |
|-----|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|---------------|
| 230 | 17.050 | 19.08 | 113.40 | 895.0 | 0.11410 | 0.15720 | 0.19100 | 0.109000 | 0.2131 |
| 529 | 12.070 | 13.44 | 77.83 | 445.2 | 0.11000 | 0.09009 | 0.03781 | 0.027980 | 0.1657 |
| 10 | 16.020 | 23.24 | 102.70 | 797.8 | 0.08206 | 0.06669 | 0.03299 | 0.033230 | 0.1528 |
| 252 | 19.730 | 19.82 | 130.70 | 1206.0 | 0.10620 | 0.18490 | 0.24170 | 0.097400 | 0.1733 |
| 46 | 8.196 | 16.84 | 51.71 | 201.9 | 0.08600 | 0.05943 | 0.01588 | 0.005917 | 0.1769 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 133 | 15.710 | 13.93 | 102.00 | 761.7 | 0.09462 | 0.09462 | 0.07135 | 0.059330 | 0.1816 |
| 509 | 15.460 | 23.95 | 103.80 | 731.3 | 0.11830 | 0.18700 | 0.20300 | 0.085200 | 0.1807 |
| 108 | 22.270 | 19.67 | 152.80 | 1509.0 | 0.13260 | 0.27680 | 0.42640 | 0.182300 | 0.2556 |
| 99 | 14.420 | 19.77 | 94.48 | 642.5 | 0.09752 | 0.11410 | 0.09388 | 0.058390 | 0.1879 |
| 145 | 11.900 | 14.65 | 78.11 | 432.8 | 0.11520 | 0.12960 | 0.03710 | 0.030030 | 0.1995 |

114 rows × 31 columns

valid



| | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | symmetry_mean |
|------------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|------------------------|---------------|
| 393 | 21.61 | 22.28 | 144.40 | 1407.0 | 0.11670 | 0.20870 | 0.28100 | 0.15620 | 0.2162 |
| 192 | 9.72 | 18.22 | 60.73 | 288.1 | 0.06950 | 0.02344 | 0.00000 | 0.00000 | 0.1653 |
| 397 | 12.80 | 17.46 | 83.05 | 508.3 | 0.08044 | 0.08895 | 0.07390 | 0.04083 | 0.1574 |
| 188 | 11.81 | 17.39 | 75.27 | 428.9 | 0.10070 | 0.05562 | 0.02353 | 0.01553 | 0.1718 |
| 437 | 14.04 | 15.98 | 89.78 | 611.2 | 0.08458 | 0.05895 | 0.03534 | 0.02944 | 0.1714 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 463 | 11.60 | 18.36 | 73.88 | 412.7 | 0.08508 | 0.05855 | 0.03367 | 0.01777 | 0.1516 |
| 545 | 13.62 | 23.23 | 87.19 | 573.2 | 0.09246 | 0.06747 | 0.02974 | 0.02443 | 0.1664 |
| 160 | 11.75 | 20.18 | 76.10 | 419.8 | 0.10890 | 0.11410 | 0.06843 | 0.03738 | 0.1993 |
| 507 | 11.06 | 17.12 | 71.25 | 366.5 | 0.11940 | 0.10710 | 0.04063 | 0.04268 | 0.1954 |
| 471 | 12.04 | 28.14 | 76.85 | 449.9 | 0.08752 | 0.06000 | 0.02367 | 0.02377 | 0.1854 |

114 rows × 31 columns

Select the features by writing the column names in the list below:

```
from imblearn.over_sampling import RandomOverSampler

def scale_dataset(dataframe, oversample=False):
    X = dataframe[dataframe.columns[:-1]].values
    y = dataframe[dataframe.columns[-1]].values

    scaler = StandardScaler()
    X = scaler.fit_transform(X)

    if oversample:
        ros = RandomOverSampler()
        X, y = ros.fit_resample(X, y)

    data = np.hstack((X, np.reshape(y, (-1, 1))))

    return data, X, y

train, X_train, y_train = scale_dataset(train, oversample=True)
valid, X_valid, y_valid = scale_dataset(valid, oversample=False)
test, X_test, y_test = scale_dataset(test, oversample=False)
```


✓ Constructing Logistic Regression Model



```
##Logistic Regression
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
y_train_pred_logistic = logistic_model.predict(X_train)
```

Display the prediction for the Logistic model against the training data (First Model)

Double-click (or enter) to edit

```
results = pd.DataFrame({'Actual': y_train, 'Predicted': y_train_pred_logistic})
results
```



| | Actual | Predicted | |
|-----|--------|-----------|---|
| 0 | 0 | 0 |  |
| 1 | 1 | 1 |  |
| 2 | 0 | 0 | |
| 3 | 0 | 0 | |
| 4 | 0 | 0 | |
| ... | ... | ... | |
| 421 | 1 | 1 | |
| 422 | 1 | 1 | |
| 423 | 1 | 1 | |
| 424 | 1 | 1 | |
| 425 | 1 | 1 | |

426 rows × 2 columns


Next steps:

[Generate code with results](#)[View recommended plots](#)[New interactive sheet](#)

Compute the accuracy of the model against the training set

```
accuracy_model1 = accuracy_score(y_train, y_train_pred_logistic)
f1_model1 = f1_score(y_train, y_train_pred_logistic)
precision_model1 = precision_score(y_train, y_train_pred_logistic)

print(f"F1 Score: {f1_model1}")
print(f"Precision: {precision_model1}")
print(f"Accuracy on training set: {accuracy_model1*100:.2f}% or {accuracy_model1}")
```



```
F1 Score: 0.9905660377358491
Precision: 0.995260663507109
Accuracy on training set: 99.06% or 0.9906103286384976
```

✓ Construct the K-Nearest Neighbors Model

```
##Import Required Libraries of KNN here
from sklearn.neighbors import KNeighborsClassifier
```


Display the prediction for the KNN model against the training data (Second Model)

```
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model = knn_model.fit(X_train, y_train)
y_train_pred_knn = knn_model.predict(X_train)
```

Compute the accuracy of the model against the training set

```
accuracy_model2 = accuracy_score(y_train, y_train_pred_knn)
f1_model2 = f1_score(y_train, y_train_pred_knn)
precision_model2 = precision_score(y_train, y_train_pred_knn)

print(f"F1 Score: {f1_model2}")
print(f"Precision: {precision_model2}")
print(f"Accuracy on training set: {accuracy_model2*100:.2f}% or {accuracy_model2}")
```

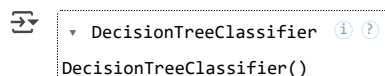


```
F1 Score: 0.9813084112149533
Precision: 0.9767441860465116
Accuracy on training set: 98.12% or 0.9812206572769953
```

✓ Constructing Decision Tree Model

```
##Decision Tree
from sklearn.tree import DecisionTreeClassifier
```

```
tree_model = DecisionTreeClassifier()
tree_model.fit(X_train, y_train)
```



DecisionTreeClassifier

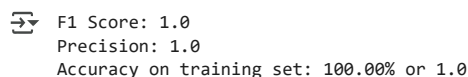
Display the prediction for the Decision Tree model against the training data (Third Model)

```
y_train_pred_tree = tree_model.predict(X_train)
```

Compute the accuracy of the model against the training set

```
accuracy_model3 = accuracy_score(y_train, y_train_pred_tree)
f1_model3 = f1_score(y_train, y_train_pred_tree)
precision_model3 = precision_score(y_train, y_train_pred_tree)

print(f"F1 Score: {f1_model3}")
print(f"Precision: {precision_model3}")
print(f"Accuracy on training set: {accuracy_model3*100:.2f}% or {accuracy_model3}")
```



```
F1 Score: 1.0
Precision: 1.0
Accuracy on training set: 100.00% or 1.0
```

✓ Results

Get the highest model for your dataset using if else condition on every f1 score / accuracy / precision of your model.

```
if f1_model1 > f1_model2 and f1_model1 > f1_model3:
    f1_best_model = "Model 1"
elif f1_model2 > f1_model1 and f1_model2 > f1_model3:
    f1_best_model = "Model 2"
else:
    f1_best_model = "Model 3"

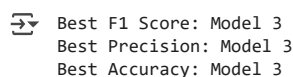
if precision_model1 > precision_model2 and precision_model1 > precision_model3:
    precision_best_model = "Model 1"
elif precision_model2 > precision_model1 and precision_model2 > precision_model3:
    precision_best_model = "Model 2"
else:
    precision_best_model = "Model 3"

if accuracy_model1 > accuracy_model2 and accuracy_model1 > accuracy_model3:
    accuracy_best_model = "Model 1"
elif accuracy_model2 > accuracy_model1 and accuracy_model2 > accuracy_model3:
    accuracy_best_model = "Model 2"
else:
    accuracy_best_model = "Model 3"
```

✓ CONCLUSION

Create conclusion and result based on the output of every model.

```
print(f'Best F1 Score: {f1_best_model}')
print(f'Best Precision: {precision_best_model}')
print(f'Best Accuracy: {accuracy_best_model}')
```



```
Best F1 Score: Model 3
Best Precision: Model 3
Best Accuracy: Model 3
```

✓ CONCLUSION

The Decision Tree model (Model 3) was found to be the best-performing classification model overall after three models were evaluated on the provided dataset: Logistic Regression (Model 1), K-Nearest Neighbors (Model 2), and Decision Tree (Model 3).

The fact that Model 3 obtained the highest results on all three evaluation metrics—F1 Score, Precision, and Accuracy—supports this conclusion. This shows that the Decision Tree model demonstrated both high recall (reducing false negatives) and good precision (reducing false positives), correctly classifying occurrences in the dataset.

The Decision Tree model finally outperformed K-Nearest Neighbors and Logistic Regression, despite their encouraging performances. To enhance these models' performance on this specific dataset, more research and adjustment may be required.

Because the Decision Tree model performs better overall in terms of F1 Score, Precision, and Accuracy, it is advised for this classification assignment. It is crucial to remember that these findings are unique to the evaluation measures and dataset in question. Conclusions may vary depending on the evaluation criteria and datasets used.

Double-click (or enter) to edit